

Technical Challenge – Portfolio Metrics Extraction

Background

Sagard receives regular reporting packages from portfolio companies across different strategies.

These reports are typically delivered as **PDFs** and include a mix of:

- Financial information (e.g., revenue, margins, cash)
- Operating metrics (e.g., ARR, churn, pipeline, headcount)
- Written commentary on performance

Each company has its own way of structuring and labelling these reports. As a result, comparing performance across companies and over time requires a lot of manual work.

We are increasingly looking to use data and software-driven approaches to reduce this manual effort and to make portfolio insights available more quickly.

Objective

Design and build a **small proof-of-concept (“crawl” phase)** that can:

- Take a folder of PDF reporting packages
- Extract a meaningful subset of metrics from those PDFs
- Organize those metrics into a structure that could be used for review and basic analysis

This is not about building a full production system. It is about how you:

- Approach an ambiguous, open-ended data problem
 - Make reasonable scoping decisions for a **1–2 hour exercise**
 - Use modern tools and techniques to work with messy documents
 - Communicate your assumptions and next steps clearly
-

Your Task

1. **Design a simple extraction approach**
 - Decide how you will read and process the PDFs.
 - Choose a small set of metrics you consider important for portfolio monitoring (for example: revenue, ARR, EBITDA, churn, headcount, or others you find meaningful).

- Make reasonable assumptions about how to recognize these metrics in the documents.
2. **Implement a lightweight pipeline or workflow**
- Build a small script, notebook, or minimal pipeline that:
 - Processes the sample PDFs
 - Extracts your chosen metrics where possible
 - Handles the fact that not every metric will appear in every PDF
 - You may use any open-source libraries or services you find appropriate, including more advanced approaches to document and text understanding if you believe they help.
3. **Organize the extracted metrics**
- Produce an output that makes it easy for a Sagard team member to review what you've extracted across companies.
 - The exact format is up to you (e.g., a table in a notebook, a CSV, a small in-memory data structure, etc.).
 - Aim for something that could naturally evolve into input for a dashboard or internal report.
4. **Capture your thinking**
- Write a short note (this can be in a README, markdown file, or within your notebook) that explains:
 - How you approached the problem
 - What you chose to implement and why
 - Key assumptions

You do **not** need to handle every edge case, every PDF perfectly, or every possible metric. We're more interested in how you navigate ambiguity and structure your solution.

Deliverables

Please submit:

- A **Github repository** with your code available to view
- The repo should include a **README** with any minimal instructions needed to run it
- Your **short written note** explaining approach, assumptions, and potential next steps

We will look through your code together in a technical conversation where you'll walk us through your decisions, tradeoffs, and ideas for how this could be extended.

Please submit a link to your repository **ahead of your interview** to allow time for review in advance.